




ECON 2041

Introductory Econometrics

Great expectations

Plan for the semester



Plan for today

- Your grade forecasts, revealed
- Income and life expectancy in the US
- Discrete or categorical?
- The correlation gallery
- Where a covariance comes from

What you expect of yourself

What I asked you

Your own semester

- Expected final mark (0-100)
- Plans:
 - attend live?
 - study hours?
 - try tutorials before class?
- The mark if you did "everything"


Three hypothetical students

- Attends live lecture:
 - rarely
 - some weeks
 - most weeks
- Study hours:
 - under 2
 - 2 to 5
 - 5+
- Try tutorials first: no / yes

For each: - a mark - % chance they pass, - % chance they get a D

Nobody expects to fail; the median is 75


Own forecast, 2026 ($N=41$), median 75




Full effort: you expect it to add about 7 marks

Own forecast, 2026 ($N=41$), median 75


Full-effort forecast, 2026 ($N=41$), median 85




The same contrast, student by student




Last year's marks: median 65, one in five failed




Attendance plans: flat, on 6 and 6 people




Study-hour plans: 72, 76, 73



Tutorials first: 70 vs 78, on 13 and 28 people




Confidence in the plan barely matters




What you expect of others


By assigned attendance: 64, 74, 76



Study hours: 64, 70, 78



Tutorials before class: 67 against 75




Independence by design


Share of hypothetical students at each study-hours level, within each attendance level

Assigned attendance	Assigned study hours per week			<i>N</i>
	Under 2	2 to 5	More than 5	
Rarely	44%	30%	26%	50
Some weeks	20%	26%	54%	35
Most weeks	42%	26%	32%	38
All vignettes	37%	28%	36%	123

Marks are not independent of attendance



Pass chance, given attendance and hours



A distinction is a pass; 10 answers disagreed

For each hypothetical student, two chances:

- $P(\text{mark} \geq 50)$, the chance of a pass
- $P(\text{mark} \geq 75)$, the chance of a distinction

Every mark of 75 or more is also a mark of 50 or more

$\{\text{mark} \geq 75\} \subset \{\text{mark} \geq 50\}$, so $P(\text{mark} \geq 75) \leq P(\text{mark} \geq 50)$

10 of 120 answers had the distinction chance above the pass chance

Income and life expectancy

Exploratory data analysis: why?

We always start by looking at the data

- Check your understanding of the variables
- Guide subsequent analysis
- Give context to results
- Generate additional questions
- Offer simple, but possibly important answers!

Descriptives can be incredibly powerful


Descriptive statistics

- mean
- median
- variance
- correlations
- conditional expectations

Life expectancy at birth

Life expectancy, 2023

The period life expectancy¹ at birth, in a given year.



Life expectancy is based on an estimate of the average age that members of a particular population group will be when they die

Source: [Our World in Data](#)

Life expectancy by current age

Life expectancy for people of different ages, Australia, 2023



The total period life expectancy¹ for people who have reached a given age.



Source: [Our World in Data](#)


Guess the gap

US men, aged 40, 2001 to 2014

How many more years does a man in the richest 1% expect to live than a man in the poorest 1%?

Portal round open now: tell your neighbor your number first, then slide

The gap is 14.6 years




Men, age 40

- Poorest 1%: 72.7 years
- Richest 1%: 87.3 years

Women: a gap of 10.1 years

Source: [Chetty et al. \(2016\)](#)

Life expectancy in the United States



Y: life expectancy at age 40

X: income percentile


$$E[Y | X = x]$$

Compare:

- $E[Y | X = p20]$
- $E[Y | X = p80]$

Source: [Chetty et al. \(2016\)](#)

Life expectancy in the United States



Y: life expectancy at age 40

X: income percentile


$$E[Y | X = x]$$

Compare:

- $E[Y | X = p20]$
- $E[Y | X = p80]$

Source: [Chetty et al. \(2016\)](#)

Life expectancy in the United States



Y: life expectancy at age 40

X: income percentile

$$E[Y | X = x, Z = z]$$


Compare:

- $E[Y | X = p20]$
- $E[Y | X = p80]$

Source: [Chetty et al. \(2016\)](#)

Unconditional vs. conditional expectation


$E[Y]$ hides a lot of heterogeneity that we can see with $E[Y|X]$!



Life expectancy, 2023



The period life expectancy¹ at birth, in a given year.

Our World
in Data




Importance of looking beyond averages!

Low-income men in the US have $\approx$ LE as Zimbabwean men!



Is this a causal relationship?




Does this suggest that giving more money to those at the bottom of the income distribution would ↑ their life expectancy?

Source: [Chetty et al. \(2016\)](#)

Descriptives can be incredibly powerful!

Place matters!



It's better to be poor in New York than in Detroit

Source: [Chetty et al. \(2016\)](#)

Discrete or categorical?

The recording's rule needs a companion

Essential Concepts, week 3: discrete takes a finite set of values, continuous takes a continuum

But every column in every dataset has finitely many values

- PISA *escs*, the socio-economic index: 62 distinct values in 12,136 rows
- PISA *math* score: 3,858 distinct values

Ask what the variable measures in the world, not how many values the column has

Three kinds of variables in the world

In the world	Example	Arithmetic makes sense?
A count	number of siblings: 0, 1, 2, ...	yes: an average of 1.4 siblings
A measurement	math score, socio-economic index, height	yes: any value in a range, rounded when recorded
A label	degree enrolled in, suburb, gender	no: there is no average suburb

Codes look numeric, and Python averages them

Education recorded as 1 = completed primary, 2 = completed high school, 3 = completed a degree

`df['education'].mean()` returns 2.3

2.3 is a fact about the codes, not about anyone's education

Rounding keeps arithmetic, binning kills it

Rounded

- A score of 61.4 recorded as 61
- Average of 61, 74, 58: 64.3, and it means something

Binned

- PISA [book](#): books at home, recorded as 0-10, 11-25, 26-100, 101-200, 201-500, 500+
- Average of "101-200" and "26-100": nothing

Books at home is a count in the world; the survey turned it into a label at collection time

Two words for the label case, both in use

- The Essential Concepts recordings say discrete for anything that is not continuous
- Categorical is the everyday word for the label kind: it names why averaging is nonsense
- You will meet both; when a variable is a label, say categorical and nobody will object


The correlation gallery

Six clouds of dots


For each one: tell your neighbor your guess, then slide to your own answer on the portal

Then we see the room's guesses against the truth


Gallery 1: math and reading, 12,136 students




Gallery 1: math and reading, $r = 0.79$




Gallery 2: math and family background




Gallery 2: math and family background, 0.38




Gallery 3: happiness and GDP, 144 countries




Gallery 3: happiness and GDP, $r = 0.69$




Same countries, log GDP: $r = 0.79$




Gallery 4: fertility and life expectancy



Gallery 4: fertility and life expectancy, -0.80




Gallery 5: the same pair, 25 students




Data: PISA 2022, Australia (OECD), 25 students drawn at random

Gallery 5: 25 students, $r = 0.49$




Data: PISA 2022, Australia (OECD), 25 students drawn at random


A thousand samples of 25: r from -0.4 to 0.8



Gallery 6: one last cloud




Gallery 6: $r = -0.03$, yet y depends on x




Where a covariance comes from

300 students, math against family background




Data: PISA 2022, Australia (OECD), 300 students drawn at random

Step 1: find the two means




Step 2: each student's two deviations




$$(x_i - \bar{x})(y_i - \bar{y}) = (+1.09)(+89) = +97$$

Step 3: the sign of the product is the quadrant



Step 4: add up: 199 pull up, 101 pull down



Nonsense units, so we standardize

$$\text{cov}(X, Y) = \frac{1}{n} \sum_i (x_i - \bar{x})(y_i - \bar{y}) = 29 \text{ index points} \times \text{score points}$$

$$r = \frac{\text{cov}(X, Y)}{s_X s_Y} = \frac{29}{0.78 \times 91.7} = 0.41$$

Some take-aways

- A conditional expectation is a function: one average for every value of X
- Averaging hides the ends: 14.6 years, inside one country
- Correlation sees straight lines only, and a sample's r is itself random
- The covariance is a vote by quadrant; dividing by the SDs removes the units